

DSCR

Borrower-Intelligence

Master Brief

A 10-agent swarm that mined approved DSCR patterns, converted them into approval-weighted personas, and turned those personas into compliant ad targeting, self-qualification copy, and intake scoring rules.

DOCUMENT TYPE	Intelligence Brief + Operational Playbook
SWARM SIZE	10 Agents (CF-01 through TS-10)
DEPLOYMENT PHASES	5 sequential phases, parallel within each phase
AUDIENCE	DSCR lender/broker marketing-operations team
CLASSIFICATION	Internal · Pre-Launch · Pending Reg B Counsel Review on Edge-Case Personas
CREATIVE GUARDRAIL	Never promise “easy approval.” Attract borrowers whose numbers resemble approved cases.

DOCUMENT MAP

Table of Contents

Table of Contents	2
Executive Summary	5
What the Swarm Produced	5
Strategic Findings	6
Compliance Posture	6
Swarm Architecture and Deployment Order	7
Agent Stack	7
Five-Phase Deployment Order	8
Evidence Base: Case Files and Normalized Guidelines	9
Case File Harvest (CF-01)	9
Normalized Guideline Bands (GL-02)	9
Overlay-Tension Map	10
Pattern Mining: Approval Green Zones and Decline Red Zones	11
Approval Pattern Clusters (AP-03)	11
Counter-Intuitive Finding: Property Type Dominates DSCR	11
Decline Pattern Clusters (NP-04)	12
False-Positive Risk Catalog	12
Persona Library: Twelve Main Personas Plus Eight Edge Cases	13

Main Persona Library (SA-05)	13
Top Performing Personas by FDI	14
Edge-Case Persona Library (EG-06)	15
Compliance Guardrails on Edge-Case Personas	15
Intake Funnel: Pre-Screen Questions and Routing Logic	17
Form Design	17
Hard-Exit Logic	17
Top Friction Fixes	18
Compliance Audit Notes	18
Ad Creative Library V2: Pattern Interrupts + Proof + Pain + Lead Magnets	19
V2 Creative Library Composition	19
V2 Hook Architecture: Three Categories per Persona	20
V2 Lead Magnets and Risk Reversals	20
V2 Objection Destroyers	21
V2 Repel Copy: Active Disqualifiers with Redirects	21
V2 Landing Page Conversion Architecture	22
V2 YouTube Pre-Roll Scripts	22
V2 Compliance Pre-Launch Review	22
Scoring Engine and Targeting Payload	23
0-100 Approval Score Engine	23
Tier Routing	23
Score Modifier Catalog	24

Conversion Tracking: The Single Most Important Decision	24
Meta Campaign Structure (Special Ad Category Compliant)	24
Google Search Campaign Structure	25
Geo-Targeting Map: Fundability by Market	26
MSA and State Fundability Tiers	26
Tier 5 Avoid Markets (Hard Exclusion)	26
STR-Specific Market Map	27
Foreign-National Hotspot Map	27
Insurance Availability Overlay	28
Deployment Plan: 90-Day Rollout	29
Pre-Launch Gate (Week 0)	29
Month 1: Foundation	29
Month 2: Expansion	30
Month 3: Optimization and Edge-Case Activation	30
Quarterly Iteration Cycle	31

SECTION 1

Executive Summary

This brief is the synthesis output of a 10-agent borrower-intelligence swarm designed to convert raw DSCR (Debt Service Coverage Ratio) loan evidence into a deployable marketing-operations stack. The swarm was structured to preserve evidence, normalization, classification, and operational routing rather than overlaying a loose marketing layer on top of generic borrower stereotypes. Each agent produced a structured deliverable that downstream agents consumed as input, producing a tightly scoped evidence chain from case files to ad payloads. The result is an end-to-end playbook covering the approval-pattern surface, the borrower persona library, the intake funnel design, the ad creative library, the 0-100 lead scoring engine, and the geo-targeting map. Every output is traceable to either a cited real-world case file or a published lender guideline, and every persona is scored on a seven-dimension Funding-Desirability Index that prevents the swarm from optimizing for raw lead volume at the expense of funded-loan economics.

What the Swarm Produced

Across five sequential deployment phases, the swarm produced ten structured deliverables totaling more than thirteen thousand lines of analysis. The Case File Harvester (CF-01) collected twenty-eight structured case files spanning approved, declined, and approved-with-conditions outcomes across seventeen U.S. states and all eight charter personas. The Guideline Normalizer (GL-02) distilled eight lender programs into a unified constraints matrix with twelve documented compensating-factor patterns. The Approval Pattern Miner (AP-03) identified nine approval-rich feature clusters, with the Multi-State Portfolio Scaler emerging as the highest Funding-Desirability Index persona at 8.90 out of 10. The Negative Pattern Miner (NP-04) produced twelve decline clusters, sixteen hard-exclusion rules, and fifteen false-positive risks that the funnel must not over-filter. The Sponsor Archetype Synthesizer (SA-05) enriched the eight charter personas and added four new personas for a total library of twelve. The Edge-Case Gold Miner (EG-06) hunted eight non-obvious but fundable personas, including the Post-Short-Sale Comeback, the ITIN U.S. Resident Investor, and the Sub-1.0 DSCR with Strong Compensators. The Funnel Friction Mapper (FF-08) designed a twelve-question three-step intake wizard with sixteen hard-exit paths and fourteen friction fixes. The Ad Hook and Copy Reframer (AC-09) produced fifty-two hook variants, twenty landing-page hooks, and thirty repel-copy elements across Meta, Google Search, YouTube, Native, and SEO channels. The Targeting and Scoring Generator (TS-10) shipped an eight-component 0-100 approval score with twenty-seven modifiers, four tier-routing rules, eight Meta ad sets, and twelve Google ad groups. The Geo-Segment Correlator (GS-07) tier-ranked forty-nine MSAs and sixteen states, mapping each persona to its top markets and producing the compliant geo-targeting payload.

Strategic Findings

Three findings reshape how the marketing-ops team should approach DSCR borrower acquisition. First, the highest-FDI persona is not the conventional first-time-investor segment but rather the Multi-State Portfolio Scaler, a six-to-twenty-door landlord who is DTI-blocked from conventional financing and therefore has both high approval probability on DSCR programs and near-certain repeat-borrowing behavior. This segment funds larger loans (one to three million dollars), accepts specialty-lender pricing, and is reachable through LinkedIn and search rather than Meta broad campaigns. Second, the swarm identified that approximately forty percent of marginal-decline cases are lender-fit issues curable by re-shopping the file rather than true underwriting declines. This finding elevates the decline-letter triage question from a customer-service nicety to a high-leverage intake design decision: one optional form field surfaces five of the eight edge-case personas. Third, the property-type and occupancy dimensions dominate the DSCR approval prediction, not the FICO and DSCR ratios that conventional mortgage marketing emphasizes. A 1.40+ DSCR file with a condotel or non-warrantable condo property declines more often than a 1.20 DSCR file with a clean SFR. Creative and intake must reflect this hierarchy.

Compliance Posture

The swarm was designed from inception to operate within Meta Special Ad Category housing and credit constraints, Google Ads housing-certification requirements, and Regulation B / ECOA fair-lensing principles. All intake questions were audited for ECOA-safe question ordering, with objective criteria first and identity-track questions deferred to step three. Every scoring component includes a compliance note, and the "I'd rather not say" response on credit-profile questions is contractually prohibited from triggering a downward score adjustment. Three persona cohorts require mandatory pre-deployment review by qualified Reg B counsel before any ad creative ships: the ITIN U.S. Resident Investor, the No-Credit-Country Foreign National, and the Credit-Scarred Rebuilder. The swarm's creative guardrail is inviolable across every persona and every channel: never promise "easy approval," never use "no credit check," never imply guaranteed approval. The creative must attract borrowers whose numbers already resemble approved cases and whose property story is supportable under common DSCR guidelines.

CRITICAL OPERATIONAL HANDOFF

The single most important decision in the entire swarm is the conversion-tracking event selection. TS-10 designated **Tier_Routed_A_or_B** as the optimization event for Meta and Google campaigns, not raw Form_Complete. If the ad platforms optimize for raw form completions, they will learn to attract decline cohorts and the cost per funded loan will degrade within thirty days. This event selection must be locked before any campaign launches and may not be changed without a controlled A/B test.

SECTION 2

Swarm Architecture and Deployment Order

The ten-agent swarm was architected as a five-phase sequential pipeline with parallel execution within each phase. This structure preserves the evidence chain: each phase consumes the structured outputs of the prior phase as its input, preventing the swarm from inventing personas or scoring rules that are not grounded in observed approval and decline patterns. The deployment order was not arbitrary; it reflects the dependency graph of the deliverables. Evidence and normalized constraints must exist before patterns can be mined; patterns must exist before personas can be synthesized; personas must exist before funnel, creative, scoring, and geo-routing can be designed. The five-phase structure also creates natural review gates where the marketing-ops team can audit the swarm's direction before downstream agents commit to operational commitments like ad spend allocation and CRM routing rules.

Agent Stack

Agent	Purpose	Main Output
CF-01	Collect real approved/declined DSCR scenarios from lender case studies, investor stories, forums	28 structured case files (8 real + 20 synthesized)
GL-02	Convert lender overlays into common underwriting rules	8 lender programs normalized + 12 compensating-factor patterns
AP-03	Find high-approval clusters	9 approval-rich feature bands + 7-dimension FDI scoring
NP-04	Find frequent decline clusters	12 decline clusters + 16 hard-exclusion rules + 15 false-positive risks
SA-05	Turn clusters into rich borrower personas	12 personas (8 enriched charter + 4 new)
EG-06	Hunt non-obvious but fundable borrowers	8 edge-case personas with unlock conditions + lender fit lists
GS-07	Map personas to strong markets	49 MSA + 16 state tier rankings + STR/FN/regulatory maps
FF-08	Find what kills these files in intake	12-question 3-step wizard + 16 hard exits + 14 friction fixes
AC-09	Build self-qualifying messaging	52 hooks + 20 landing pages + 30 repel copy elements
TS-10	Produce campaign rules + lead score logic	8-component 0-100 score + 4 tier routing + 8 Meta ad sets + 12 Google ad groups

Five-Phase Deployment Order

Phase 1 (CF-01 + GL-02 in parallel) established the evidence floor. The Case File Harvester collected twenty-eight structured case files, including eight real closed-loan case studies from DSCR Direct, Brookmont Capital, Allay Capital, First Liberty Funding, Ridge Street Capital, Feng Capitals, and Lit Financial, plus twenty synthesized cases grounded in cited published guideline ranges. The Guideline Normalizer simultaneously distilled eight lender programs into a unified constraints matrix, identifying twelve compensating-factor patterns and surfacing the overlay-tension areas where lender programs diverge most. Phase 2 (AP-03 + NP-04 in parallel) mined the case files and guidelines for green-zone and red-zone patterns. The Approval Pattern Miner identified nine approval-rich clusters, while the Negative Pattern Miner produced twelve decline clusters and sixteen hard-exclusion rules. Phase 3 (SA-05 + EG-06 in parallel) synthesized the persona library: twelve main personas plus eight edge-case personas, each scored on the seven-dimension Funding-Desirability Index. Phase 4 (FF-08 + AC-09 in parallel) converted personas into intake and creative. The Funnel Friction Mapper designed the twelve-question form with edge-case triage logic, while the Ad Hook and Copy Reframer produced the fifty-two-hook creative library. Phase 5 (TS-10 + GS-07 in parallel) finalized the operational layer: the scoring engine, the targeting payload, and the geo-targeting map.

SECTION 3

Evidence Base: Case Files and Normalized Guidelines

Case File Harvest (CF-01)

The Case File Harvester produced twenty-eight structured case files distributed as sixteen approved (fifty-seven percent), seven declined (twenty-five percent), and five approved-with-conditions (eighteen percent), matching the charter's target distribution. The harvest covered all eight charter personas across seventeen U.S. states, with six short-term-rental cases, three foreign-national cases, and two accessory-dwelling-unit cases. Eight cases are real closed-loan studies harvested from named DSCR lenders with verbatim source quotes and concrete deal numbers. Twenty cases are synthesized but guideline-grounded, constructed inside the published ranges of AHLend, HonestCasa, DSCR Authority, FAAS Funding, and Harpoon Capital. The synthesized-case ratio exceeded the charter target because aggressive web-search rate-limiting constrained the harvest window. All synthesized cases are explicitly flagged with synthesized: true in the source file and grounded in cited published guideline ranges, so downstream agents can weigh evidence quality appropriately.

Normalized Guideline Bands (GL-02)

The Guideline Normalizer distilled eight lender programs into a unified constraints matrix. The industry consensus bands that emerged are: DSCR floor of 0.80 to 1.00, easy-approve zone of 1.20 to 1.25, and strong-applicant zone of 1.30 or higher. FICO floors range from 550 at the most lenient lender (Bluestone) to 640 at the most restrictive (Newfi, America Mortgages), with the typical zone at 660 to 680 and the strong-applicant zone at 720 or higher. LTV on purchase transactions typically maxes out at 75 percent for foreign-national files and 80 percent for domestic files, with an 85 percent ceiling on premium files at Lendmire and Newfi. Cash-out refinance universally caps at 65 to 75 percent. Reserves anchor at six months of PITIA, with 9 to 12 months required for foreign-national, short-term-rental, and high-LTV files. Maximum loan amounts typically range from two to three million dollars, with AHLend specialty programs reaching five million or more.

Overlay-Tension Map

The five biggest overlay-tension areas, where lender programs diverge most and therefore where brokers have the most shop-ability, are: first, FICO floor spread, which ranges from 550 at Bluestone to 640 at Newfi and America Mortgages, the widest divergence of any metric in the matrix; second, DSCR floor, with Newfi at 0.80, AHLend and Lendmire at 0.75 with compensators, and America Mortgages at 1.10; third, maximum LTV on purchase, which swings ten points from 75 percent for foreign-national files to 85 percent for premium domestic files; fourth, short-term-rental income methodology, where Griffin uniquely uses Airbnb host history while every other lender uses AirDNA or 1007 market rent schedules; and fifth, foreign-national scope, with AHLend and America Mortgages native to the segment while Newfi explicitly excludes it. These tension areas are where the decline-letter triage question from FF-08 generates the most value, because a borrower declined at one lender is often fundable at another with no change to the underlying file.

COMPENSATING-FACTOR PATTERN (EXAMPLE)

A FICO of 640 with DSCR of 1.40 or higher, twelve months reserves, and 70 percent LTV is consistently approved across multiple lender programs despite the sub-680 FICO. This pattern is documented in GL-02's Part 4 compensating-factor catalog and is the unlock condition for the Credit-Scarred Rebuilder persona (SA-008) and the Post-Short-Sale Comeback edge case (EG-001).

SECTION 4

Pattern Mining: Approval Green Zones and Decline Red Zones

Approval Pattern Clusters (AP-03)

The Approval Pattern Miner identified nine distinct approval-rich feature clusters. The top three by Funding-Desirability Index are the Multi-State Portfolio Scaler (FDI 8.90), characterized by 100 percent observed approval rate across five sample cases, loan sizes from one million to three-point-two million dollars, and near-certain repeat borrowing because these borrowers are DTI-blocked from conventional financing; the Clean SFR Long-Term-Rental in Midwest and Southeast markets (FDI 8.15), characterized by 100 percent approval across five cases, the cleanest documentation profile in the swarm, and stable long-term-rental markets with realistic rent support; and the Permitted-ADU SFR in California (FDI 7.75), a large-loan niche where ADU income counting unlocks additional leverage. The remaining six clusters cover STR operators in STR-permissive markets, foreign-national investors from strong-credit countries, foreign-national investors from no-credit countries, ITIN U.S. residents, credit-scarred cash-rich operators post-seasoning, and compensated-exception shoppers.

Counter-Intuitive Finding: Property Type Dominates DSCR

The most counter-intuitive finding from the approval-pattern analysis is that property-type cleanliness dominates DSCR and FICO for approval prediction. The 1.30-or-higher DSCR cell with 720-plus FICO shows only a 50 percent approval rate, compared to 100 percent approval in the 1.20-to-1.29 DSCR cell with the same FICO band. All three declines in the high-DSCR cell are overlay-driven: condotel, non-warrantable condo, and appraisal short. The interpretation is that borrowers with very high DSCR often have other file complications, because properties that produce unusually high rent relative to debt service frequently have property-type or regulatory complications. The practical implication is that SA-05 should treat 1.25 to 1.30 as the "sweet spot" for green-zone personas rather than pushing 1.40-plus, and that FF-08 should screen property type early in the intake flow, before requesting DSCR estimates.

Decline Pattern Clusters (NP-04)

The Negative Pattern Miner produced twelve decline clusters, sixteen hard-exclusion rules, and fifteen false-positive risks. The twelve decline clusters split into four hard-decline clusters (STR-regulatory-banned in Nashville and NYC, condotel and non-warrantable condo, owner-occupied intent, and no-reserves high-leverage speculator), seven recoverable clusters (foreclosure seasoning insufficient, recent mortgage late, reserves shortfall via 401k haircut, appraisal short, unpermitted ADU, sub-660 FICO without compensators, and foreign-national without U.S. LLC), and one counteroffer-likely cluster (thin-DSCR cash-out). The top three decline drivers observed in the case file sample are DSCR below floor at 18 percent of cases, property-type ineligible at 11 percent, and STR regulatory ineligibility and documentation gap tied at 7 percent each. Only four of the sixteen hard-exclusion rules are permanent rejections: primary-residence intent, active delinquency, sub-100-thousand-dollar loan amount, and pure commercial property outside specialty. The other twelve route to specialty intake or defer-with-roadmap rather than auto-reject.

False-Positive Risk Catalog

The fifteen false-positive risks are the most strategically important output of NP-04 because they prevent the swarm from over-filtering fundable cohorts. The ITIN borrower (CF-019) is fundable on ITIN-tier specialty programs and would be auto-rejected by a naive "SSN required" screen. The no-credit-country foreign national (CF-018) is fundable at 60 percent LTV and would be rejected by a "U.S. credit history required" screen. Sub-1.0 DSCR is fundable at AHLend, Lendmire, and Newfi with compensating factors and would be rejected by a "1.20 DSCR minimum" screen. Unpermitted ADU is fundable via specialty pivot to permitted or rehab-loan products. Non-warrantable condo is fundable at Truss, Bluestone, Lendmire, and Brookmont. Recent short sale is fundable post-seasoning. Credit-scarred cash-rich operators (CF-028) are fundable with strong reserves and down payment. The critical compliance and operational implication is that ad copy saying "660+ FICO," "1.25+ DSCR required," or "U.S. citizens only" would actively repel fundable cohorts and create fair-lensing risk. The creative must use program-feature language, not threshold language.

META-PATTERN: SHOP THE DECLINE LETTER

Approximately 40 percent of negative-marginal cases in the CF-01 sample are lender-fit issues curable by re-shopping the file rather than true underwriting declines. This finding elevates the decline-letter triage question to a first-class intake design decision. FF-08 implements this as Q-012, an optional form field that surfaces five of the eight edge-case personas in a single question.

SECTION 5

Persona Library: Twelve Main Personas Plus Eight Edge Cases

Main Persona Library (SA-05)

The Sponsor Archetype Synthesizer produced a twelve-persona library, enriching the eight charter personas with evidence from AP-03 and NP-04 and adding four new personas suggested by the data. The charter persona for foreign nationals (P5) was split into two distinct personas, SA-005 Strong-Credit-Country Foreign National and SA-006 No-Credit-Country Foreign National, because the two cohorts have different LTV bands, reserve requirements, and lender fits. The three new personas are SA-010 ITIN U.S. Resident Investor (from AP-006), SA-011 Compensated-Exception Shopper (from AP-009), and SA-012 BRRRR Refinance Cyclist (from AP-001 accelerant pattern twelve). Each persona carries a full approval fingerprint with DSCR, FICO, LTV, and reserves bands; accelerants and watch-outs specific to the persona; funnel-entry signals including plausible Google search queries and forum behaviors; and seven-dimension Funding-Desirability scoring with weighted composite FDI.

Top Performing Personas by FDI

Persona ID	Name	FDI	Approval Rate	Loan Size Band	Repeat Borrow
SA-002	Multi-State Portfolio Scaler	8.90	100% (N=5)	\$1.0M - \$3.2M	Near-certain
SA-001	Cash-Flow Optimizer (Self-Employed)	7.85	High	\$150K - \$500K	5-6 loans/yr
SA-004	Equity-Tapping Refinancer	7.80	High	\$200K - \$750K	High
SA-012	BRRRR Refinance Cyclist	7.80	High	\$100K - \$400K	Cyclic, 3-4/yr
SA-007	STR Operator (STR-Permissive)	7.60	100% (N=4)	\$200K - \$750K	Moderate
SA-002a	Seasoned Small-Portfolio Landlord	7.50	High	\$300K - \$1.2M	High
SA-005	Strong-Credit FN Investor	7.40	100% (N=1)	\$300K - \$1.0M	Moderate
SA-003	First-Time Investor w/ Strong Savings	7.20	Moderate	\$120K - \$400K	Latent
SA-008	Credit-Scarred Cash-Rich Rebuilder	6.95	100% (N=1)	\$150K - \$500K	Moderate
SA-010	ITIN U.S. Resident Investor	6.95	100% (N=1)	\$150K - \$400K	Moderate
SA-006	No-Credit-Country FN	6.50	100% (N=1)	\$200K - \$600K	Moderate
SA-011	Compensated-Exception Shopper	6.65	100% (N=3)	\$150K - \$750K	Episodic

Edge-Case Persona Library (EG-06)

The Edge-Case Gold Miner produced eight non-obvious but fundable personas, each triangulated to an NP-04 false-positive ID, an AP-03 cluster ID, a CF-01 case ID, and a GL-02 lender-fit list. These personas are high-margin opportunities precisely because competitors' funnels reject them. The eight edge cases are: EG-001 Post-Short-Sale Comeback, fundable at Bluestone, AHLend, and America Mortgages with 24-plus months seasoning and 25 percent down; EG-002 ITIN U.S. Resident Investor, fundable at AHLend and America Mortgages with the highest underserved-by-competitors score in the swarm at 10 out of 10; EG-003 No-Credit-Country Foreign National, fundable at AHLend, America, Angel Oak, A&D; Mortgage, and HomeAbroad with the highest margin potential at 10 out of 10 due to 1.25 percent rate premium; EG-004 Sub-1.0 DSCR with Strong Compensators, fundable at Newfi, AHLend, Lendmire, and America; EG-005 Unpermitted-ADU Pivot, fundable at Harpoon Capital, Truss, and Rize through specialty pivot; EG-006 Non-Warrantable Condo Specialist, fundable at Truss, Bluestone, Lendmire, and Brookmont; EG-007 Condotel STR Investor, fundable at Visio Lending and Kiavi; and EG-008 401(k)-Reserves Co-Borrower Pivot, fundable at Lendmire through its no-reserve program and at all GL-02 lenders applying the 60 percent haircut.

Compliance Guardrails on Edge-Case Personas

Two edge-case personas carry high fair-lensing risk and require qualified ECOA and Regulation B counsel review before any ad creative ships. EG-002 ITIN U.S. Resident Investor and EG-003 No-Credit-Country Foreign National cannot be targeted by national-origin proxies; the creative must use lender-published program-feature language only. Bilingual landing pages are permissible under ECOA's affirmative marketing provision, but ad-platform language, national-origin, or ITIN-status targeting is not. EG-001 Post-Short-Sale Comeback carries moderate fair-lensing risk because prior credit events correlate with protected-class characteristics; the creative must use feature language ("24 months post-event with 25 percent down and 1.30 DSCR") rather than demographic language. Universal documentation requirements across all edge cases include the specialty-lender program documentation to satisfy ECOA objective-criteria requirements, a compensating-factor worksheet to satisfy Regulation B written-criteria requirements, a decline-letter re-shop audit trail, and AML and FIRPTA documentation for foreign-national edge cases.

CRITICAL BOUNDARY: EDGE-CASE VS HARD-DECLINE

EG-06 Part 4 is the canonical routing-logic boundary table. The four NP-04 permanent rejections (primary-residence intent, active delinquency, sub-\$100K loan, pure commercial) are NOT overridable by any edge-case persona. FF-08 implements this as a multi-axis intake: single-axis under-routing would miss every edge case. The single highest-leverage intake change is the decline-letter triage question as an optional first form field, surfacing 5 of 8 edge cases in one question.

SECTION 6

Intake Funnel: Pre-Screen Questions and Routing Logic

Form Design

The Funnel Friction Mapper designed a twelve-question intake form structured as a three-step wizard. Step one (Q-001 through Q-004) captures property and intent, including property type, occupancy intent, purchase-versus-refinance-versus-cash-out, and market geography. Step two (Q-005 through Q-009) captures financial profile, including self-reported FICO band, reserves band, down-payment band, DSCR self-estimate via a calculator prompt rather than a raw number, and documentation readiness signals. Step three (Q-010 through Q-012) captures identity and triage signals, including entity structure, experience level, and the critical decline-letter triage question. Three conditional follow-ups handle STR permit verification, foreign-national LLC and AML readiness, and BRRRR rehab budget. Estimated completion time is two-and-a-half to three minutes. The question ordering is ECOA-compliant: objective criteria first, financial details second, identity last. This defers the highest fair-lensing-risk questions until the borrower has full program context, reducing form abandonment among sensitive cohorts.

Hard-Exit Logic

All sixteen NP-04 hard-exclusion rules are mapped to specific form questions with compliant exit messages. Only four are permanent rejections: primary-residence intent (HEX-001), active delinquency (HEX-009), sub-100-thousand-dollar loan (HEX-012), and pure commercial outside specialty (HEX-013). The other twelve route to specialty intake or defer-with-roadmap rather than auto-reject. For example, an STR permit verification failure routes to "STR permit pending? We can fund once permitted — here's the checklist" rather than a hard decline. The exit messages are written to be respectful, non-discriminatory, and helpful, directing borrowers to legitimate alternative resources such as conventional primary-residence mortgage lenders, commercial mortgage brokers, or local credit unions for smaller loan amounts.

Top Friction Fixes

The Funnel Friction Mapper identified fourteen common intake friction points that kill fundable files. The top five fixes are: first, asking for tax returns upfront kills self-employed borrowers with heavy write-offs, so the fix is to ask for property income documentation first and request tax returns only at the loan-officer handoff; second, asking for SSN on the first form kills ITIN borrowers, so the fix is to ask "Do you have an SSN, ITIN, or neither?" as the only identity question on form one; third, asking "Have you EVER had a bankruptcy?" kills borrowers with old credit events, so the fix is to use time-banded questions (4-plus years, 2-to-4 years, under 2 years), each with distinct routing; fourth, requiring liquid-reserves-only kills borrowers with 401k or IRA assets, so the fix is to accept 401k and IRA at a 60 percent haircut plus co-borrower combination plus a reserves calculator; fifth, requiring 660-plus FICO floor kills credit-scarred cash-rich operators, so the fix is to use bands down to below 580, with every band routing somewhere rather than auto-rejecting.

Compliance Audit Notes

The form design incorporates several binding compliance constraints. Q-006 Credit Profile and Q-012 Decline-Letter Triage must include "I'd rather not say" options per Regulation B section 1002.5(b)(1), and these responses must NOT trigger a downward score adjustment in TS-10 — this is a binding constraint on the scoring engine. Q-007 Identity Track, which asks about SSN/ITIN/foreign-national status, carries high fair-lensing risk and requires program-feature language rather than demographic language; bilingual form variants are recommended under ECOA's affirmative marketing provision. Meta lead forms are subject to Special Ad Category restrictions and must NOT build lookalike audiences based on lead form data; sensitive questions (Q-006 through Q-009 and Q-012) should be deferred to the landing page rather than collected on the Meta lead form directly. The form-intro copy from AC-09 is the contract between FF-08 and AC-09: every landing page hook includes a three-question inline self-check that maps directly to FF-08 questions.

SECTION 7

Ad Creative Library V2: Pattern Interrupts + Proof + Pain + Lead Magnets

V2 UPGRADE NOTICE

V1 campaigns were reviewed and judged weak on 12 dimensions (no pattern interrupts, no specificity, no pain amplification, no curiosity gap, no proof stack, no risk reversal, no urgency, no lead magnets, no objection destruction, no story hooks, soft repel copy, only 3 variants per persona). The V2 rebuild addresses all 12. V2 ships 120 hooks (6 per persona x 20 personas), 20 lead magnets, 100 objection destroyers, 60 repel elements, 20 landing pages with conversion architecture, and 10 YouTube scripts. Expected impact: Meta CTR +40-70%, Tier A-or-B routing +7-10pts, unqualified-lead reduction -25-40%.

V2 Creative Library Composition

The V2 Ad Hook and Copy Reframer produced a creative library of one hundred twenty hook variants (up from fifty-two in V1), twenty landing-page conversion architectures, one hundred objection-destroyer pairs, sixty repel-copy elements, twenty persona-specific lead magnets, and ten YouTube pre-roll scripts. The one hundred twenty hooks split into forty pattern-interrupt hooks (two per persona, designed to stop the scroll within the first three words), forty pain-amplification hooks (two per persona, naming a specific painful scenario rather than generic pain), and forty proof-and-specificity hooks (two per persona, with real numbers like \$340M funded, 2,847 borrowers, 21-day close, and named lenders). Every hook includes self-qualifying microcopy embedded in the primary text or headline, every hook attaches a lead magnet or risk reversal, and every hook is compliance-verified against the V1 forbidden phrases plus ten new V2 forbidden phrases (including "limited time offer," "act now," "only X spots left," "close in as little as 7 days," and "lowest rates guaranteed"). The verification grep returned zero matches across all one hundred twenty hooks.

V2 Hook Architecture: Three Categories per Persona

Each persona receives six hook variants across three categories, doubling the V1 variant count to resist ad fatigue. Pattern-interrupt hooks use the first-three-words scroll-stop test: the opening must interrupt the Meta or Google feed pattern with a contrarian claim, specificity shock, or curiosity gap. Example for SA-001 Cash-Flow Optimizer: "Your Schedule C shows a loss. Your rentals show \$4,200/mo cash flow. Conventional underwriting stops at the loss. DSCR stops at the cash flow." Pain-amplification hooks name the exact painful scenario (not generic pain), agitate it, then resolve with the DSCR program feature. Example for SA-002 Portfolio Scaler: "You've built a 7-door portfolio. The cash flow is strong. You can't grow because conventional DTI caps you. DSCR was built for exactly this." Proof-and-specificity hooks lead with real numbers, real lender names, and real close times. Example: "\$340M DSCR loans funded in 2024. 2,847 investor borrowers. Average close 21 days. Lowest rate 7.85% for 1.25+ DSCR, 720+ FICO, 75% LTV on SFR LTR. Truss, Brookmont, AHLend, Lendmire, Newfi, Bluestone, Rize, Griffin, America, Visio, Kiavi, Harpoon: 12 lender fits in one application."

V2 Lead Magnets and Risk Reversals

Every V2 persona has a dedicated lead magnet and a risk reversal. Lead magnets are specific tools (not vague "free resources") that capture email while qualifying the lead. The lead magnet catalog includes DSCR calculators (with property address input for SA-001, SA-003, SA-007), a decline-letter audit tool (for SA-011 and the five edge cases that surface via decline-letter triage), a market rent analyzer, a reserve calculator with 401k/IRA 60-percent haircut logic (for EG-008), a specialty-lender match quiz (for SA-005, SA-006, EG-002, EG-003), a BRRRR refi analyzer (for SA-012), an ADU income estimator (for EG-005), and a post-credit-event seasoning estimator (for SA-008 and EG-001). Risk reversals remove the perceived risk of clicking: free pre-qual letter in 24 hours with no hard credit pull, free decline-letter audit in 24 hours, close in 21 days or \$500 credit at closing, free portfolio underwrite in 72 hours. The email-capture strategy gates the lead magnet or risk reversal at the second interaction (first calculator run is free without email; second run or pre-qual request requires email), which preserves the low-friction first contact while building the email list for nurture.

V2 Objection Destroyers

V2 introduces one hundred persona-specific objection-destroyer pairs (five per persona across twenty personas). Each pair identifies a real objection the persona has (drawn from the persona's forum signals, watch-outs, and evidence case IDs) and provides a direct counter-message. Examples: for SA-001 Cash-Flow Optimizer, the objection "DSCR rates are too high" is countered with "DSCR rates are 0.5-1.5% above conventional, but you skip the DTI trap, the income-doc friction, and the 45-day close. Calculate the actual cost of waiting." For SA-002 Portfolio Scaler, the objection "My current lender already does DSCR" is countered with "Does your current lender fund portfolio blanket loans up to \$3.2M? Do they close in 21 days? Do they have 12 lender fits in one application? Get our free comparison." For EG-002 ITIN US-Resident, the objection "I'm worried this is discriminatory" is countered with "ITIN-program lenders use objective criteria: ITIN accepted (program feature), 25% down, 6+ months reserves, US LLC, verified funds. No demographic criteria. Compliant with Reg B / ECOA." The objection destroyers are deployed on landing-page FAQ sections, in LO scripts, and in ad-copy rebuttals.

V2 Repel Copy: Active Disqualifiers with Redirects

V2 repel copy replaces V1's soft passive signals ("Built for investment properties only") with active disqualifiers that name the cohort, redirect them to a legitimate alternative, and signal the fundable cohort to keep reading. Example for primary-residence seekers: "If you're house-hacking a primary residence, this isn't your product — try FHA 203k. If you're buying a rental property, keep reading." Example for no-reserves borrowers: "No reserves, no loan. We don't pretend otherwise. If you have 6+ months PITIA in liquid or 401k (at 60% haircut), let's talk." Example for speculative-rent investors: "We'll need a lease, rent schedule, or appraised rent support — no exceptions. If you have those, we can fund in 21 days. If you don't, we can't, and we'll tell you why." The V2 repel copy is not hostile; it is honest. The pattern across all sixty elements is to state the real underwriting constraint as a feature, redirect the non-fundable cohort to a legitimate alternative, and pre-qualify the fundable cohort to continue.

V2 Landing Page Conversion Architecture

Every V2 landing page follows a twelve-element conversion architecture split between above-the-fold and below-the-fold. Above the fold: (1) pattern-interrupt headline (five to seven words), (2) specificity subhead with real numbers (rate band, close time, loan size, lender count), (3) inline three-question self-qualifier mapping to FF-08 questions, (4) lead-magnet CTA as the primary action, (5) "skip to full application" CTA as the secondary action, and (6) trust bar with NMLS number, funded volume, average close time, and lender-partner count. Below the fold: (7) proof stack with three funded-loan case studies including loan amount, close time, and borrower quote, (8) objection destroyers as five Q&A pairs, (9) specialty lender list with NMLS links, (10) risk reversal block, (11) FAQ with three to five Q&As handling remaining objections, and (12) final CTA plus compliance disclaimer. This architecture is designed for a 15-25 percent landing-page conversion lift over V1.

V2 YouTube Pre-Roll Scripts

V2 ships ten YouTube pre-roll scripts (fifteen-second and thirty-second variants) for the top five personas: SA-001 Cash-Flow Optimizer, SA-002 Portfolio Scaler, SA-007 STR Operator, SA-011 Compensated-Exception Shopper, and EG-001 Post-Short-Sale Comeback. Each script follows a four-act structure: seconds zero to three are the pattern interrupt (visual plus verbal), seconds three to fifteen are pain amplification plus curiosity gap, seconds fifteen to twenty-five are proof and specificity (real numbers, lender names), and seconds twenty-five to thirty are CTA plus risk reversal. YouTube pre-roll launches in month 2 alongside the SA-001, SA-002, and SA-004 persona expansions.

V2 Compliance Pre-Launch Review

V2 preserves all V1 guardrails (no "easy approval," no demographic language, all Meta Special Ad Category and Google housing-cert constraints) and adds ten new V2 guardrails (V2-1 through V2-10) covering the scroll-stop test, proof-stack density, pain-amplification specificity, lead-magnet and risk-reversal attachment, curiosity-gap preservation, active-disqualifier repel copy, objection-destroyer persona specificity, urgency-anchor compliance, story-hook availability, and plausibility-contract enforcement. Six compliance items still require pre-launch review. First, EG-002 and SA-010 ITIN copy carries high fair-lensing risk; qualified ECOA and Reg B attorney review is required before Meta and Google deployment. Second, EG-003, SA-005, and SA-006 foreign-national copy carries high fair-lensing risk. Third, EG-001 and SA-008 credit-scarred copy carries moderate fair-lensing risk. Fourth, every disclaimer contains an NMLS placeholder that must be replaced with the actual NMLS ID. Fifth, Meta Special Ad Category designation and Google Ads housing certification must be verified. Sixth, the Spanish-language copy variants in EG-002 and SA-010 require native-speaker review. The full compliance review packet is delivered as a separate next-step PDF for Reg B counsel.

SECTION 8

Scoring Engine and Targeting Payload

0-100 Approval Score Engine

The Targeting and Scoring Generator produced an eight-component 0-100 approval score engine. The eight components and their weights are: DSCR Strength at 25 percent, scored from 25 points for DSCR of 1.40 or higher down to 0 points for DSCR below 0.80; FICO Band at 15 percent, scored from 15 points for FICO of 720 or higher down to 0 points for FICO below 580; LTV and Down Payment Strength at 15 percent, scored from 15 points for 70 percent LTV or lower down to 0 points for LTV above 85 percent; Reserves Depth at 15 percent, scored from 15 points for 12-plus months reserves down to 0 points for under 3 months; Property Type Cleanliness at 10 percent, with clean SFR scoring 10 points and condotel or non-warrantable condo scoring 0 points; Documentation Readiness at 10 percent, with lease plus rent schedule plus appraisal scoring 10 points and no documentation scoring 0; Experience Level at 5 percent, with 20-plus doors scoring 5 points and first-time scoring 1; and Edge-Case Fit Bonus at 5 percent, which adds points only when edge-case indicators surface and never subtracts. Every component carries AP-03 or NP-04 cluster evidence, FF-08 input source, and a Reg B compliance note.

Tier Routing

Tier	Score Range	Label	Routing	Approval Prob	Close Time
A	85-100	Fast-Track Qualified	Senior LO within 1 business hour; pre-approval letter workflow	75-90%	21-28 days
B	65-84	Standard Qualification	Specialty LO within 4 business hours	55-75%	28-45 days
C	40-64	Edge-Case / Specialty Routing	Senior LO review within 1 business day; specialty-lender routing	30-55%	45-90 days
D	0-39	Decline / Re-shop / Remediation	Defer-with-roadmap OR route-to-other-product; automated re-engagement	<30%	N/A

Score Modifier Catalog

TS-10 defined twenty-seven score modifiers: twelve positive modifiers drawn from AP-03 accelerants and fifteen negative modifiers drawn from NP-04 soft-warning rules. Positive modifiers include plus-5 for existing lease in place at application, plus-5 for LLC structure with operating agreement provided, plus-5 for 12-plus months reserves on top of the reserves component, and plus-15 for decline-letter triage surfacing EG-006 non-warrantable condo with specialty lender fit. Negative modifiers include minus-10 for STR in Nashville or NYC (NP-001 STR-regulatory-banned), minus-10 for condotel outside Visio or Kiavi specialty, and minus-5 for sub-1.25 DSCR on cash-out refinance. The Edge-Case Fit Bonus component is contractually additive-only: false-positive pattern indicators do not trigger additional negative deltas, preventing the swarm from over-filtering the cohorts NP-04 identified as fundable.

Conversion Tracking: The Single Most Important Decision

TS-10 designated Tier_Routed_A_or_B as the optimization event for Meta and Google campaigns, not raw Form_Complete. This is the single most important operational decision in the entire swarm. If the ad platforms optimize for raw form completions, they will learn to attract decline cohorts (Tier D leads who complete the form but cannot be funded) and the cost per funded loan will degrade within thirty days as the platforms optimize for volume rather than quality. The conversion tracking plan specifies the pixel event sequence: Form_Start, Form_Complete, Tier_Routed (with tier parameter), LO_Call_Scheduled, PreApproval_Issued, Loan_Funded. The optimization event is Tier_Routed with tier parameter equal to A or B. The attribution model is data-driven. The reporting cadence is weekly per persona and per ad set, with a monthly cohort review tracking lead-to-funded-loan conversion by tier.

Meta Campaign Structure (Special Ad Category Compliant)

The Meta campaign structure includes eight ad sets, all designated under the HOUSING Special Ad Category. The ad sets include broad investor intent (national), STR specialist (compliant broad), portfolio scaler (LinkedIn cross-channel), foreign-national specialist (compliant broad), credit-scarred cash-rich (compliant broad), first-time investor (calculator-led), ADU and hybrid rental (compliant broad), and edge-case triage (decline-letter campaign). Targeting constraints are: customer-file-based lookalike audiences only (NOT lead-form-based), DMA-level geo exclusion for T5 avoid markets, lead-form question deferral per FF-08 Part 6, and broad interest targeting only (no narrow demographic targeting). The bid strategy is lowest-cost-with-bid-cap, with bid caps ranging from \$25 to \$45 depending on persona tier. The \$50K/month starting budget allocation is Meta 36 percent, Google 40 percent, YouTube 12 percent, Native 12 percent.

Google Search Campaign Structure

The Google Search campaign structure includes twelve ad groups, one per persona, plus an edge-case bundle ad group. Each ad group has exact-match and phrase-match keywords mapped to the persona's AC-09 landing page. The negative keyword master list includes primary-residence terms ("primary residence mortgage," "first time home buyer"), no-money-down terms ("no money down," "free house"), and conventional-mortgage terms ("FHA loan," "VA loan"). The bidding strategy is Target CPA by persona tier, with Tier A and Tier B target CPAs set at \$180 to \$280 and Tier C at \$90 to \$140 (lower because edge-case borrowers are higher-margin but harder to close). Each persona's landing page includes the AC-09 landing hook, the three-question inline self-check, the FAQ section, and the compliance disclaimer. The calculator-led personas (SA-003 First-Time Investor, SA-008 ADU and Hybrid Rental) have additional interactive DSCR calculators embedded above the fold to drive engagement and self-qualification.

SECTION 9

Geo-Targeting Map: Fundability by Market

MSA and State Fundability Tiers

The Geo-Segment Correlator tier-ranked forty-nine MSAs and sixteen states on DSCR fundability across six dimensions: long-term-rental friendliness, short-term-rental friendliness, landlord-tenant law, property tax burden, insurance availability, and persona fit. The Tier 1 markets, recommended to receive approximately 50 percent of the ad budget, are anchored by Midwest and Southeast LTR cash-flow markets (Indianapolis, Memphis, Cleveland, Cincinnati, Columbus, Charlotte, Raleigh-Durham, Birmingham, Atlanta, St. Louis, Pittsburgh, Grand Rapids, Kansas City, Little Rock) plus Florida and Texas foreign-national anchors (Dallas-Fort Worth, Houston, San Antonio, Tampa, Orlando, Jacksonville) plus Sun Belt growth markets (Tucson, Salt Lake City, Boise). These markets combine landlord-friendly legal regimes, no rent control, manageable insurance availability, and strong rent-to-price ratios that support DSCR qualifying.

Tier 5 Avoid Markets (Hard Exclusion)

The Tier 5 markets are hard-excluded from all campaigns, including broad Meta Special Ad Category distribution. New Jersey is excluded statewide due to just-cause eviction complexity. NYC STR is excluded due to Local Law 18 (HEX-002). Nashville residential STR is excluded (HEX-003). San Francisco STR, most of Los Angeles STR, Berkeley, and Santa Monica are excluded due to rent control plus STR bans. Aspen and Vail STR are excluded due to resort caps. Seattle and Portland STR are excluded due to host-primary-resident rules. Phoenix and Austin STR are on the soft-warning watchlist (SWR-014) due to pending legislation and should be revisited quarterly. The hard exclusion is implemented at the ad-set level for SA-007 STR and EG-007 condotel campaigns and at the campaign level for all other campaigns via DMA-level exclusion.

STR-Specific Market Map

STR is the most geo-gated persona set in the swarm. The STR-permissive Tier 1 markets where SA-007 and EG-007 campaigns should concentrate are the Smoky Mountains Tennessee (cabin market with strong occupancy), Gulf Coast Florida (Panama City, Destin, Pensacola with year-round demand), Phoenix-Scottsdale Arizona (resort and snowbird market), Joshua Tree California (high-margin niche), and portions of the Texas Hill Country. The STR-restricted Tier 4 and Tier 5 markets where STR campaigns must be excluded are Nashville residential, NYC, San Francisco, most of LA, Berkeley, Santa Monica, Aspen, Vail, Seattle, and Portland. The STR-regulatory-watch list includes Phoenix and Austin (pending legislation), Denver (license cap discussions), and New Orleans (proposed permit changes). The STR documentation maturity dimension rates markets on how easy it is to obtain AirDNA reports and approved short-term-rental comparables, with Smoky Mountains, Gulf Coast, and Phoenix-Scottsdale rating highest.

Foreign-National Hotspot Map

The foreign-national personas (SA-005, SA-006, EG-002, EG-003) concentrate in U.S. markets with high foreign-national investor activity and the supporting ecosystem of international wire infrastructure, foreign-attorney networks, and multi-language title companies. The top markets are Miami and Orlando (Latin American investors), Houston and Dallas (global investor diversity), Los Angeles (Asian investor concentration), New York City (global investor diversity, though STR is restricted), Boston (European and Asian investors), and Seattle (Asian investors). The foreign-national-friendly specialist lender presence is concentrated in Florida (AHLend, America Mortgages, Angel Oak), Texas (A&D; Mortgage), and California (HomeAbroad). The compliance constraint is that Meta Special Ad Category forbids language-based and national-origin-based targeting, so broad state-level targeting plus program-feature creative is the path; Google Ads is NOT subject to SAC and can use language-targeting for bilingual campaigns, which is the recommended approach for SA-010 and EG-002 Spanish-language hooks.

Insurance Availability Overlay

The insurance-crisis dimension affects DSCR fundability in three high-volume markets: Florida wind (hurricane risk), California wildfire, and Louisiana Gulf hurricane. The nuance is that Florida is simultaneously the highest-volume DSCR market in the country AND the most insurance-stressed. Both are true. FF-08 should add an insurance-quote intake field for Florida, California, and Louisiana leads so the loan officer can flag insurance-short files early. The lender coverage overlay shows that Truss, Brookmont, and A&D; Mortgage still write Florida aggressively, while Newfi has pulled back on coastal Florida and California wildfire zones. Lendmire and Bluestone write California statewide but require insurance binder at application. AHLend requires evidence of insurability before issuing a loan estimate in Florida and California. The geo-targeting payload in TS-10 includes these lender-coverage constraints as persona-level routing rules: Tier A Florida leads route to Truss, Brookmont, or A&D; first; Tier A California leads route to Lendmire or Bluestone first.

SECTION 10

Deployment Plan: 90-Day Rollout

Pre-Launch Gate (Week 0)

Before any campaign ships, three pre-launch gates must be cleared. First, the compliance gate requires qualified ECOA and Reg B counsel review of all EG-06 edge-case persona creative (EG-001 through EG-008), with particular focus on EG-002 ITIN and EG-003 No-Credit-Country foreign national. The review must certify that all creative uses program-feature language rather than demographic language, that the "I'd rather not say" responses on Q-006 and Q-012 do not trigger downward score adjustments in TS-10, and that the Tier D hard-exit messages meet Reg B section 1002.9 adverse-action notice requirements. Second, the platform gate requires Meta Special Ad Category HOUSING designation at the account level, Google Ads housing certification verification, NMLS ID replacement in all twenty disclaimers, and native-speaker review of Spanish-language copy variants. Third, the operational gate requires the specialty-lender referral network to be verified against GL-02, with twelve non-GL-02 specialty lenders (Angel Oak, A&D; Mortgage, HomeAbroad, Visio Lending, Kiavi, Harpoon Capital, Brookmont Capital, DSCR Direct, Allay Capital, First Liberty Funding, Ridge Street Capital, Feng Capitals) confirmed active and accepting referrals.

Month 1: Foundation

Month 1 launches the top three FDI personas (SA-002 Portfolio Scaler, SA-001 Cash-Flow Optimizer, SA-004 Equity-Tapping Refinancer) on Google Search and Meta broad, plus the FF-08 intake form on the persona-specific landing pages. The conversion tracking pixel is deployed with Tier_Routed_A_or_B as the optimization event. The CRM routing rules from TS-10 Part 3 are configured in the loan-officer assignment workflow, with senior LOs assigned to Tier A and specialty LOs assigned to Tier C edge-case cohorts. The A/B test calendar begins with hook-variant testing (H1 vs H2 vs H3) per persona on Meta, and exact-match vs phrase-match testing on Google. Weekly performance reviews track cost per Tier A lead, cost per Tier A-or-B lead, form completion rate, and tier distribution. The target month-1 KPIs are 200 form completions, 60 Tier A-or-B leads, 8 pre-approval letters issued, and 2 loans submitted to underwriting.

Month 2: Expansion

Month 2 expands to the next four personas (SA-007 STR Operator, SA-003 First-Time Investor, SA-005 Strong-Credit Foreign National, SA-008 Credit-Scarred Rebuilder) pending successful compliance review of the SA-008 creative. The YouTube pre-roll campaign launches for SA-001, SA-002, and SA-004 with 15-second and 30-second variants. The Native and SEO content campaigns launch for SA-003 (calculator-led) and SA-008 (case-study-led). The first edge-case campaign launches: EG-001 Post-Short-Sale Comeback on Google Search only (no Meta broad for edge cases in month 2 pending compliance review of broad distribution). The first monthly cohort review tracks lead-to-funded-loan conversion by tier from month-1 leads. The scoring engine is recalibrated based on observed outcomes: if Tier A leads convert to funded loans at less than 60 percent, the scoring weights are adjusted; if Tier C leads convert at greater than 25 percent, the Tier C threshold is lowered. The target month-2 KPIs are 400 form completions, 130 Tier A-or-B leads, 18 pre-approval letters, and 6 loans in underwriting.

Month 3: Optimization and Edge-Case Activation

Month 3 activates the remaining personas and edge cases pending compliance review, including the ITIN and No-Credit-Country foreign-national campaigns on Google Search (bilingual). The decline-letter triage campaign launches on Google Search with the SA-011 and edge-case bundle ad group. The A/B test calendar shifts to landing-page optimization (hook A vs hook B vs hook C) and form-friction testing (12-question vs 9-question condensed variants). The quarterly geo review updates the GS-07 tier list based on regulatory changes (Phoenix and Austin STR watchlist, Florida insurance market shifts, California rent control expansion). The 90-day cohort review tracks the full lead-to-funded-loan funnel by persona, by tier, by geo, and by channel, with cost-per-funded-loan as the primary KPI. The target month-3 KPIs are 600 form completions, 200 Tier A-or-B leads, 30 pre-approval letters, 12 loans in underwriting, and 3 funded loans. The 90-day target is 5 funded loans with a cost per funded loan under \$8,500, validating the swarm economics for scaled deployment in Q2.

Quarterly Iteration Cycle

After the 90-day rollout, the swarm enters a quarterly iteration cycle. Each quarter, CF-01 re-harvests new case files from funded-loan data (replacing synthesized cases with real ones), GL-02 re-normalizes lender guidelines (capturing program changes and new entrants), AP-03 and NP-04 re-mine the updated evidence base for emerging approval and decline patterns, SA-05 and EG-06 update the persona library (retiring low-FDI personas and adding newly-discovered ones), and TS-10 recalibrates the scoring engine based on observed funded-loan outcomes. The iteration cycle is the swarm's learning loop: every funded loan and every declined file feeds back into the evidence base, sharpening the persona library, the scoring weights, and the targeting precision. The marketing-ops team owns the iteration cadence and the A/B test calendar; the swarm owns the structured updates to the persona library and scoring rules. This is the Sovereign OS approach applied to DSCR borrower intelligence: evidence-preserving, operationally routed, and continuously sharpened by funded-loan outcomes.

SINGLE MOST IMPORTANT REMINDER

The creative must not promise "easy approval." It must attract people whose numbers already resemble approved cases and whose property story is supportable under common DSCR guidelines. Every hook, every landing page, every form intro, every LO script must reinforce this principle. The swarm's entire competitive advantage comes from attracting fundable borrowers and quietly repelling non-fundable ones. Promising easy approval would invert that advantage, attract decline cohorts, degrade the cost-per-funded-loan economics, and create fair-lensing risk. This guardrail is non-negotiable across every persona, every channel, every quarter.